THE RATTY2TAILS PAPERS

MONETARY HISTORY SERIES

Yield Curve Control

What happens when a central bank stops letting the bond market set its own price — and what happens, later, when the market stops believing it.

EDITED BY

Ratty2Tails

SEPTEMBER 2026

A long-form essay on pegging the price of government debt

TL;DR

- **Yield curve control** is a central bank naming a price, not a quantity: it commits to buy whatever amount of a given bond is needed to hold its yield at a chosen level, rather than committing to buy a fixed amount and letting the yield fall where it will.
- The tool has been tried four times in the last ninety years: the United States pegged Treasury yields through the Second World War and Korea (1942–1951), Japan ran the longest peacetime version on record (2016–2024), Australia tried a shorter pandemic-era version that collapsed in weeks (2020–2021), and the Federal Reserve seriously studied and then declined to adopt it in 2020.
- Every episode tells the same story from a different angle: a yield peg is cheap to defend exactly as long as the market believes the central bank will defend it, and ruinously expensive — or simply impossible — the moment the market decides to test that belief.
- Japan sustained its peg for nearly eight years; Australia's collapsed inside a few weeks once inflation surprised to the upside. The difference was not the mechanism. It was credibility.
- With U.S. long-bond yields near two-decade highs, questions about the Federal Reserve's independence back in the headlines, and debt-to-GDP levels that once justified the postwar peg reappearing, yield curve control is no longer a historical curiosity — it is a live line item in the current policy debate.

INTRODUCTION & SCOPE

Naming a Price Instead of a Quantity

Most of what a central bank does to a bond market, it does by quantity. Quantitative easing announces a sum — so many billions of bonds, purchased over so many months — and lets the price that emerges from those purchases be whatever the market decides it should be. Yield curve control inverts that logic entirely. Instead of naming an amount and accepting whatever yield results, the central bank names a yield and commits to buying whatever amount is required to hold it there, for as long as it says it will. The quantity becomes the open variable; the price becomes the promise.

That inversion sounds like a technical footnote. It is not. A central bank that targets quantity always knows, in advance, the maximum size of its own commitment. A central bank that targets price does not — it has, in effect, written the market a blank cheque, redeemable at a fixed exchange rate, for as long as the peg holds. This essay is about the four times in the last ninety years that a major central bank has written that cheque: the United States during and after the Second World War, Japan for the better part of a decade, Australia for eighteen months during the pandemic, and the Federal Reserve in 2020, which studied the idea in detail and chose, deliberately, not to sign it.

The four episodes are not a random sample. Read together, they are closer to a controlled experiment in what actually holds a yield peg up. The mechanism is identical in every case — unlimited purchases at a fixed price. What varies is everything around it: the state of inflation, the depth of the domestic buyer base, the length of time the market is asked to believe the promise, and, above all, whether the market ever actually tests it. That last variable turns out to matter more than any other, which is the argument this essay is built to make.

SECTION I

Pegging a Number Instead of Buying a Pile

Mechanically, yield curve control is executed one of two ways. A central bank can run a fixed-rate, unlimited-quantity purchase operation: it announces it will buy a specific bond, or maturity bucket, at a specific price, in whatever amount the market wants to sell at that price. Or it can run repeated, open-ended interventions whenever the market yield drifts away from a stated target, buying as much as it takes to bring the yield back. Japan used the first method almost from the start; Australia and the wartime United States leaned closer to the second. Both versions share the defining feature: the central bank has fixed the price and left the quantity of its own balance sheet open-ended.

The appeal of doing this, rather than running ordinary quantitative easing, is mostly about efficiency and communication. A credible peg does not actually require the central bank to buy very much, most of the time — if the market believes the peg will hold, arbitrage does most of the work for free, and traders sell the pegged bond back toward the target rather than testing a level the central bank has promised to defend. Australia's experience makes the point vividly: the Reserve Bank spent eighteen months barely needing to intervene, because the target itself, once established, became the market price. A quantity-based programme offers no equivalent shortcut; the central bank has to keep buying the announced amount whether or not anyone believes the policy will work.

The price of that efficiency is that the central bank has traded a known cost for an unknown one. Under quantitative easing, the size of the balance-sheet expansion is chosen in advance and cannot exceed it. Under yield curve control, the balance-sheet expansion is whatever the market decides to sell the central bank at the pegged price — a number that is small when the peg is credible and can become arbitrarily large, in principle, the moment it is not. Every case study that follows is really a study of that one exposure: what happens to a central bank that has promised a price without capping the bill.

SECTION II

The Original Peg

The first major yield curve control programme belonged to the country least associated with the idea today. In April 1942, four months after the United States entered the Second World War, the Federal Reserve formally committed, at the Treasury's request, to hold short-term Treasury bill yields at three-eighths of one percent and to implicitly cap long-term Treasury bond yields at 2.5 percent. The goal was straightforward: the government was about to borrow an enormous sum to fight a war, and it wanted to borrow it cheaply, predictably, and without the market's yields wandering upward as new debt was issued.

To hold the peg, the Fed gave up control of its own balance sheet in exactly the way the mechanism implies. It bought whatever quantity of Treasuries the market offered at the pegged price, which meant its portfolio size and composition were dictated by investor behaviour rather than by any monetary-policy judgement of its own. For as long as wartime price controls kept inflation in check, this cost very little to sustain. The trouble began once the war, and the price controls that came with it, ended.

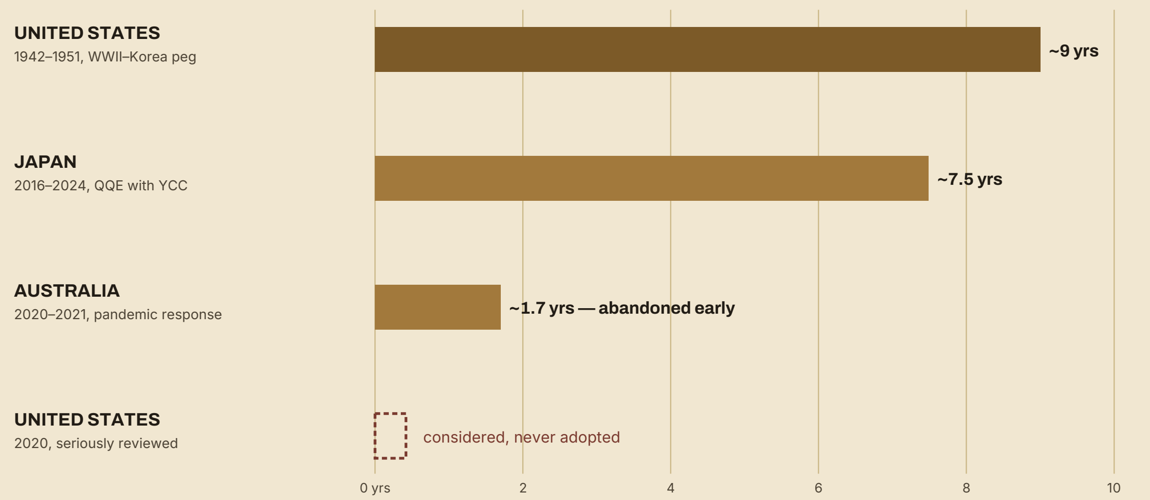

FIG. 1 — Four attempts to control the yield curve, by duration. Sources: Federal Reserve History; Bank of Japan; Reserve Bank of Australia; FOMC minutes, 2020.

Postwar inflation, and then the outbreak of the Korean War in June 1950, put the peg under real pressure for the first time. With the Treasury still insisting the low-rate commitment continue — both to keep the government's own borrowing costs down and, as the Truman administration argued at the time, to protect the value of the war bonds millions of households were holding — the Federal Open Market Committee found itself trapped by its own promise. Inflation was running above eight percent year-over-year by early 1951, and every FOMC member understood that defending a 2.5 percent bond yield in that environment meant buying an ever-larger share of new Treasury issuance directly, which meant creating an ever-larger amount of new money to do it. The central bank had become, in substance if not in name, the government's financing arm.

> *"The Treasury and the Federal Reserve System have reached full accord… to assure the successful financing of the Government's requirements and, at the same time, to minimize monetization of the public debt."*
>
> TREASURY–FEDERAL RESERVE ACCORD, MARCH 1951

The March 1951 Accord that ended the peg is remembered today, correctly, as the founding document of the modern independent Federal Reserve. What is less often noted is how narrowly it arrived: the Accord followed nearly a year of open dispute between the Fed and the Truman administration, leaked FOMC minutes, and a direct meeting between the full Committee and the President himself. The peg did not end because the Fed decided, calmly, that independence was the better policy. It ended because holding the peg any longer had become incompatible with controlling inflation at all — the first and clearest illustration of a pattern that recurs in every subsequent episode this essay covers.

SECTION III

Japan's Long Run

No central bank has run a peacetime yield curve control programme for longer than the Bank of Japan. Launched in September 2016 as part of a framework the bank called "Quantitative and Qualitative Monetary Easing with Yield Curve Control," the policy set an explicit target of around zero percent for the ten-year Japanese government bond, defended through unlimited fixed-rate purchase operations, alongside a short-term policy rate of minus 0.1 percent. The stated goal was to fight two decades of deflation and stagnant growth by anchoring long-term borrowing costs near zero indefinitely, giving households and firms every reason to expect that cheap credit was not a temporary condition but a standing one.

What made the Japanese version durable where others were not was not the mechanism — it was identical in principle to the American wartime peg and the Australian pandemic-era target — but the environment it operated in. Japan entered the policy already deflationary, with a domestic investor base heavily and voluntarily concentrated in JGBs, and with global interest rates near zero across most of the advanced world for much of the programme's life. The peg was, for years, simply not costly to defend, because there was little market pressure pushing against it in the first place.

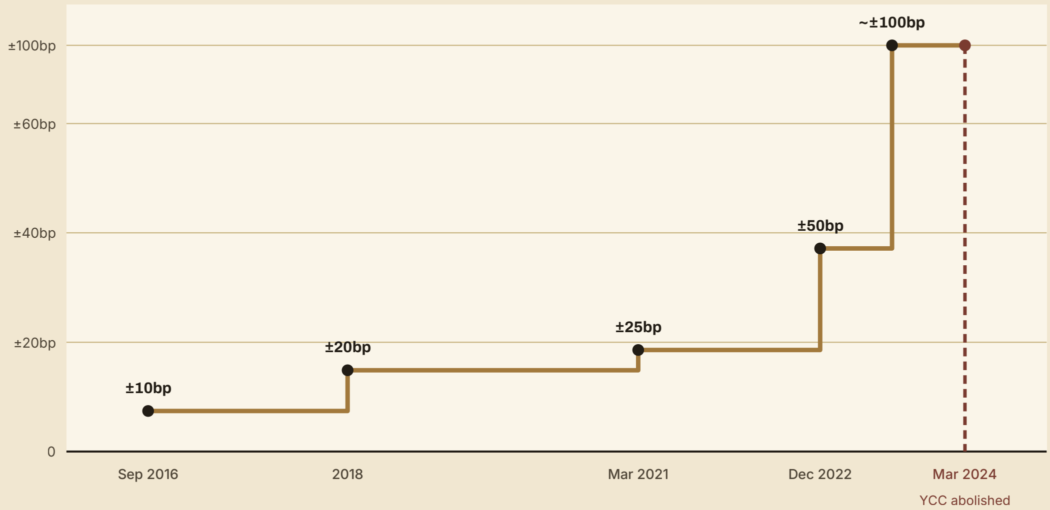

FIG. 2 — The Bank of Japan's tolerance band around its 10-year JGB yield target, widened in five discrete steps before the policy's abolition. Source: Bank of Japan policy statements, 2016–2024.

That did not mean the policy stood still. As the figure above shows, the Bank of Japan widened its tolerance band around the zero target five times over the programme's life — from roughly ten basis points either side in 2016, out to around fifty basis points by December 2022, and effectively to about one percent by July 2023, when the bank recharacterised its cap as a "reference" rather than a rigid ceiling. Each widening was, in substance, a small, managed retreat — an acknowledgement that holding the exact original line was becoming harder to justify as inflation, after decades of absence, finally showed up. The programme ended altogether on 19 March 2024, alongside the end of negative interest rates, once above-target inflation and unusually strong wage settlements gave the bank enough political cover to let the peg go without it reading as a policy failure.

Eight years is, by the standard of every other episode in this essay, an extraordinary run. But it is worth being precise about what sustained it: not a demonstration that yield curve control is inherently stable, but a nine-year stretch in which the policy was rarely seriously tested, followed by an orderly retreat once it finally was. The next case study shows what happens when the test arrives all at once.

SECTION IV

The Australian Cautionary Tale

In March 2020, as the pandemic hit, the Reserve Bank of Australia announced a target of around 0.25 percent for the yield on the three-year Australian government bond, defended through bond purchases, as an extension of its existing overnight cash-rate target further out along the curve. It took the bank just eleven days and A$27 billion in purchases to establish the target credibly. In November 2020, as the cash rate itself was cut, the RBA lowered the bond-yield target to around 0.10 percent. For the next eighteen months, the policy worked almost exactly as the theory predicts a credible peg should: the market simply adopted the RBA's target as the going rate, and the bank barely had to intervene at all.

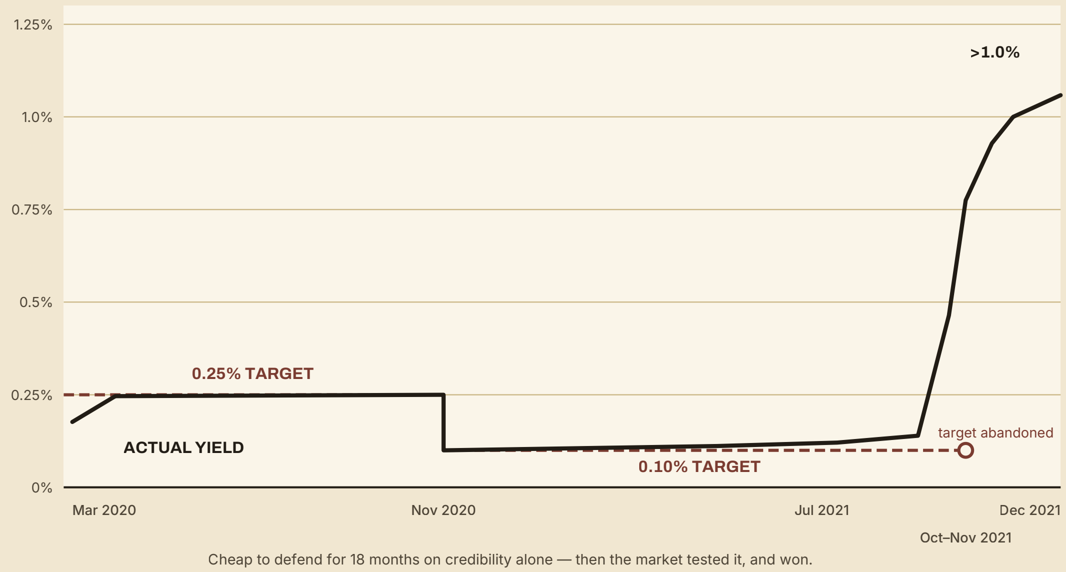

FIG. 3 — The RBA's 3-year bond yield target versus the actual market yield, March 2020–December 2021. Sources: Reserve Bank of Australia; J.P. Morgan Asset Management; Bloomberg.

Then, in October 2021, it stopped working. Stronger-than-expected employment data, a jump in core inflation back inside the RBA's target band for the first time in six years, and a broader global repricing of central-bank rate paths combined to convince bond traders that the RBA's 0.10 percent target, and its accompanying promise of no rate rises before 2024, no longer matched where policy was actually headed. On 15 October, the three-year yield began drifting above target. The RBA bought bonds to defend it — and then, within days, stopped. By the end of the month the yield had closed above 0.77 percent; within another fortnight it had traded above 1.2 percent, more than ten times the level the bank had spent eighteen months defending. On 2 November 2021, the RBA formally abandoned the target altogether, two and a half years earlier than the policy's original horizon implied.

The Reserve Bank's own subsequent review of the episode is unusually candid for a central-bank document. It describes the exit as "disorderly," acknowledges the affair caused the bank "reputational damage," and accepts that defending the target in its final weeks came at a real financial cost as yields moved against it. What the review does not describe, because it does not need to, is any change in the mechanism itself between the eighteen quiet months and the two violent weeks. The tool was identical throughout. What changed was whether the market believed the RBA would use it.

A NOTE ON CREDIBILITY

Japan and Australia ran, in essence, the same policy. One held for the better part of a decade; the other broke in a fortnight. The difference was not resolve, or purchasing power, or technical design — the RBA had, if anything, an easier target to defend, a three-year bond rather than a ten-year one. The difference was that Australia's macro backdrop turned against the peg suddenly and all at once, while Japan's turned only gradually, giving its central bank room to widen the band in five small steps rather than defend an untenable line all the way to a cliff edge.

SECTION V

The Road Not Taken

The Federal Reserve came close to adopting yield curve control once more, in the middle of 2020, and its decision not to is instructive precisely because of how carefully it was made. With the policy rate already near zero and the pandemic recession deepening, senior officials floated the idea publicly — New York Fed president John Williams said in May 2020 that the Committee was "thinking very hard about" the tool, and Vice Chair Richard Clarida confirmed later that year that yield caps remained in the Fed's toolkit. Behind the scenes, FOMC staff conducted a formal review of the three episodes then available as precedent: the American wartime peg, the Bank of Japan's programme, and the Reserve Bank of Australia's newly launched target, with the Australian case judged the most directly relevant to the Fed's own circumstances.

The review concluded against adoption. Minutes from the Committee's June and July 2020 meetings record that most participants judged the likely benefits modest, on the grounds that the Fed's forward guidance about the future path of rates was already highly credible on its own, leaving little for an explicit yield target to add. Against that modest upside, participants weighed two concrete costs: the risk of an excessively rapid expansion of the balance sheet if the peg were tested, and the practical difficulty of designing and communicating an exit — precisely the problem that would visibly undo the Reserve Bank of Australia a year later. The Committee chose to keep the tool on the shelf rather than in use, while explicitly noting it "should remain an option... if circumstances changed markedly."

That is a genuinely rare thing in the history of unconventional monetary policy: a major central bank examining a tool in real time, in detail, against live precedent, and declining to use it — not out of principle, but because its own analysis said the trade was not worth making yet.

The caveat attached to that decision is the part worth remembering. The Fed did not say the tool was flawed. It said the tool was not needed *at that time*, and left the door open for a time when it might be.

SECTION VI

Whether the Circumstances Have Changed

By the second half of 2026, several of the conditions the 2020 FOMC discussion treated as reasons not to act have moved in the opposite direction. Long-dated U.S. Treasury yields have climbed to levels not sustained since before the 2008 financial crisis, with the thirty-year yield trading near five percent and the ten-year near 4.7 percent through the summer — a rise driven partly by the scale of new issuance needed to fund persistent federal deficits, and partly by a rising term premium as investors demand more compensation for holding long-duration government debt at all. Debt-to-GDP has moved back above the range that once made the postwar peg necessary in the first place.

Questions about the Federal Reserve's independence, which the 2020 review could take largely for granted, have also become an active part of the market conversation. The Trump administration's attempt to remove Fed Governor Lisa Cook was blocked by a district court and is set to be reviewed by the Supreme Court; more broadly, sustained public pressure from the administration for faster rate cuts has fed investor concern that monetary policy could come to be shaped by the government's own financing needs rather than by the Fed's inflation mandate alone. Former Federal Reserve Chair Janet Yellen used a January 2026 conference address to name the risk directly, warning that fiscal dominance would likely raise term premiums and borrowing costs as investors grow concerned the government will lean on inflation or financial repression to manage its debt. Federal Reserve Chair Jerome Powell has, for his part, publicly rejected the suggestion that debt-servicing costs influence the Committee's rate decisions.

None of this means yield curve control is imminent, and this essay takes no view on whether it should be adopted; the question is genuinely contested among economists and carries real trade-offs whichever way a central bank resolves it. What the 2020 episode does establish is the specific, named condition under which the Fed itself said it would reconsider: a marked change in circumstances. Rising long-term yields, a widening fiscal deficit, and open public debate about central-bank independence are precisely the kind of circumstances the 2020 minutes had in mind. Several market research desks were, by late 2025 and into 2026, already sketching yield curve control as one plausible late stage in a longer sequence — a shift toward tolerating higher inflation, followed, if unaddressed, by outright caps on long-term yields. Whether that sequence plays out is a question this series will keep returning to. What is no longer true is that the question is purely academic.

SUMMARY

A Promise, Not a Purchase

Yield curve control is not, at bottom, a purchasing programme. It is a promise about price, backed by a purchasing programme large enough to keep the promise credible. Every episode this essay has traced — the American wartime peg, Japan's eight-year run, Australia's eighteen quiet months and violent unwinding, and the Fed's deliberate 2020 refusal — is a variation on that single fact. The policy is nearly free when the market believes it, and becomes either unaffordable or simply undeliverable the moment the market decides to test it, because a central bank that has pegged a price rather than a quantity has no natural limit on what defending that price might eventually cost. The tool has returned to the policy conversation in the United States not because anyone has forgotten how Australia's version ended, but because the conditions that make the promise harder to keep — large deficits, rising term premiums, and open questions about who a central bank ultimately answers to — are the same conditions every prior episode was built on.

REFERENCES

Sources & Further Reading

1. **Federal Reserve Bank of Richmond / Federal Reserve History.** “The Treasury-Fed Accord,” covering the 1942 rate peg and the March 1951 Accord.
2. **Board of Governors of the Federal Reserve System.** Minutes of the Federal Open Market Committee, 9–10 June 2020 and 28–29 July 2020 — internal review of yield caps and targets (YCT).
3. **Bank of Japan.** “Quantitative and Qualitative Monetary Easing with Yield Curve Control,” policy statement, 21 September 2016; subsequent band-widening statements, March 2021, December 2022, and July 2023; policy statement ending YCC and negative rates, 19 March 2024.
4. **Reserve Bank of Australia.** “Review of the Bond Yield Target,” 2022, covering the March 2020 introduction, the November 2020 and 2021 target changes, and the November 2021 discontinuation.
5. **J.P. Morgan Asset Management.** “The Reserve Bank of Australia’s Crisis of Confidence,” 4 November 2021.
6. **Yellen, J.** Remarks on Federal Reserve independence and fiscal dominance, Allied Social Science Associations annual meeting, 6 January 2026.
7. **Reinhart, C. M. & Sbrancia, M. B.** (2011, rev. 2015). “The Liquidation of Government Debt.” NBER Working Paper 16893; IMF Working Paper WP/15/7.
8. **U.S. Treasury market data, 2026.** Long-bond and ten-year Treasury yield levels, as widely reported through the second and third quarters of 2026.

Ratty2Tails
Editor, The Ratty2Tails Papers — Monetary History Series.
Typeset in Fraunces, Playfair Display, Source Serif 4, EB Garamond, Archivo, and Inter.